ECE2013 / Fall 2026

Planning and Games

Constraint Satisfaction and Adversarial Search

Assignment 2 / Group 6

Course Advanced AI Algorithms and Systems

Members Wang Libin Jiang Jiahao

Wang Zizheng Xue Xiaokang

Section I Constraint Satisfaction

Exact-cover formulation

The tiling problem is represented as exact cover. Each available board cell and each tile identity contributes one constraint. A legal placement covers the cells it occupies and the identity of its tile. Selecting placements that satisfy every constraint exactly once therefore covers the board without overlap or omission and uses each tile once.

Let C be the available cells and I the tile identities, giving the disjoint constraint set $U = C \sqcup I$. A legal placement p covers a subset $R_p \subseteq U$, and the binary variable x_p indicates whether it is selected. The tiling conditions become

$$\sum_{p:\, u \in R_p} x_p = 1 \quad (\forall u \in U), \qquad x_p \in \{0, 1\}. \tag{1}$$

Cell constraints enforce coverage without overlap; identity constraints enforce single use. Identical shapes still need separate identities. For example, sixty cells and twelve pentominoes give seventy-two constraints; each placement covers six: five cells and one identity.

Algorithm X and minimum remaining values

Exact cover can be solved by Algorithm X, a depth-first backtracking procedure.[1]

We first generate distinct rotations and reflections and enumerate placements that remain inside the board and avoid holes. Search uses the minimum remaining values heuristic: choose the uncovered constraint with the fewest available placements. After selecting a placement, remove all conflicting candidates; backtrack as soon as any constraint has no candidate. Integer bitsets reduce the cost of these set operations. Every candidate for the chosen constraint is considered, so the search is complete, although its worst-case complexity remains exponential.

If A denotes the available placements and U' the uncovered constraints, the minimum remaining values rule selects

$$c^* = \arg\min_{c \in U'} \big|\{p \in A : c \in R_p\}\big|. \tag{2}$$

After choosing a placement p covering c^*, remove the constraints in R_p and all other placements that intersect R_p. Zero candidates means failure; no uncovered constraints means success. Choosing the tightest constraint exposes failure early and allows search to switch adaptively between placing a tile and covering a cell.

Selected placements remain pairwise compatible, while uncovered constraints must remain coverable by subsequent candidates. Removing conflicts performs forward checking; backtracking restores the previous candidate state. The heuristic changes exploration order, not the solution set. Bitsets reduce per-node work without changing worst-case complexity.

[1] D. E. Knuth, *Dancing Links*, 2000. https://arxiv.org/abs/cs/0011047. Our candidate sets use integer bitsets, not the paper's doubly linked list implementation.

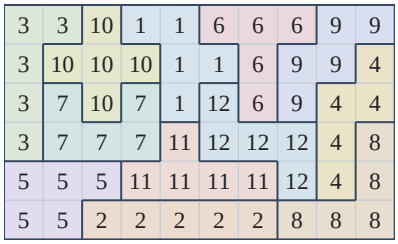
Results

Four boards were tested with dominoes, triominoes, and pentominoes. All twelve instances passed the supplied correctness checker. The single-run times in Table 1 include placement generation and search.

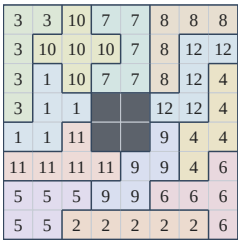
Table 1: Tiling results across boards and tile sets.

Board	Tile size	Placements	Search nodes	Time / ms
6×10	2	3120	31	14.74
6×10	3	3600	40	18.40
6×10	5	2056	863	17.01
5×12	2	3090	31	14.05
5×12	3	3520	27	17.75
5×12	5	1936	1073	16.71
3×20	2	2910	31	13.64
3×20	3	3040	31	15.46
3×20	5	1236	238	9.59
8×8 (central hole)	2	3000	31	14.70
8×8 (central hole)	3	3280	28	23.40
8×8 (central hole)	5	1568	2104	19.63

The number of placements alone does not predict search effort. For example, the 6×10 pentomino instance has fewer placements than its domino counterpart but requires more backtracking. The interactions between tile shapes matter more than the candidate count alone.



(a) 6×10 board



(b) Board with a central hole

Figure 1: Pentomino tilings. Equal labels identify cells belonging to one tile.

Section II Predefined Agents

Rules and search settings

The game ends when either player first completes a line on any local board. Local boards are indexed from 0 to 8 in row-major order. A move's position within its local board determines the opponent's next local board. If that board is full, the opponent may choose any empty global cell. All four games start on board 4. Offensive X moves first in games 1–2; defensive O moves first in games 3–4.

The predefined evaluation has three ordered tiers. Terminal states receive 10000 for an X win and -10000 for an O win. Otherwise, X scores 500 per unblocked two-in-a-row and 100 per blocked O two-in-a-row. O scores -100 per unblocked two-in-a-row and -500 per blocked X two-in-a-row. Only when the second tier has no counts are occupied local corners scored: 30 each for X and -30 each for O. Rows, columns, and diagonals are counted separately, so a cell may contribute to more than one line.

Both algorithms search three plies and stop early at terminal states. Leaves use the root agent's evaluation throughout a search, always expressed as X utility: X maximizes and O minimizes. Moves are visited in row-major order, and the current best move is replaced only by a strict improvement. Thus ties retain the first move. Alpha–beta search prunes when $\alpha \geq \beta$.

Minimax recurrence and pruning

A state must include the global board B, the forced local board k, and the player p: $s = (B, k, p)$. Piece locations alone do not determine legal actions. A move at global coordinates (r, c) sets $k' = 3(r \bmod 3) + (c \bmod 3)$.

Let $\mathcal{A}(s)$ be the legal actions, $T(s, a)$ the successor, and d the remaining depth. For a fixed root evaluation E,

$$V_d(s) = \begin{cases} u(s), & s \text{ is terminal}, \\ E(s), & d = 0 \text{ and } s \text{ is nonterminal}, \\ \max_{a \in \mathcal{A}(s)} V_{d-1}(T(s,a)), & p = X, \\ \min_{a \in \mathcal{A}(s)} V_{d-1}(T(s,a)), & p = O. \end{cases} \tag{3}$$

Here $u(s)$ is -10000, zero, or 10000 for an X loss, draw, or win. Minimax assumes the opponent chooses the action most adverse to this utility. The resulting conservative decision is optimal for the truncated tree, not necessarily for the full game, since nonterminal E is approximate.

Alpha and beta record existing MAX guarantees and MIN restrictions along the path. Suppose MAX already has an option worth 5. If another MIN branch finds a child worth 3, that branch's final value cannot exceed 3, so its remaining children cannot improve MAX's choice. Pruning preserves the minimax value for the same depth and evaluation.[2] For uniform branching factor b and depth d, ordinary search costs $O(b^d)$. With ideal move ordering, pruning can approach $O(b^{d/2})$; its worst case remains $O(b^d)$. The required row-major order does not justify assuming ideal ordering in these experiments.

[2] D. E. Knuth and R. W. Moore, *An Analysis of Alpha-Beta Pruning*, Artificial Intelligence 6 (1975), 293–326. https://doi.org/10.1016/0004-3702(75)90019-3.

Outcomes and pruning

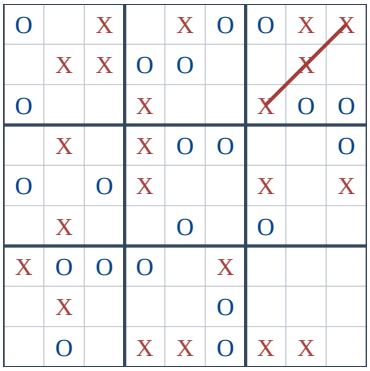
Node counts include visited roots and leaves and are summed over all turns. In Table 2, MM denotes minimax and AB denotes alpha–beta.

Table 2: Results of the four required matchups.

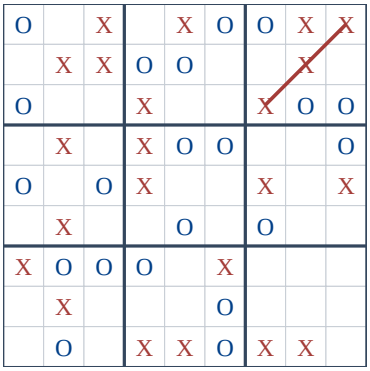
Game	X	O	First	Plies	X nodes	O nodes	Total	Win
1	MM	MM	X	43	8768	7948	16716	X
2	MM	AB	X	43	8768	3398	12166	X
3	AB	MM	O	44	3353	8528	11881	X
4	AB	AB	O	44	3353	3431	6784	X

X won all four games. Games 1–2 have identical move sequences, with a 27.2% reduction in visited nodes after pruning. Games 3–4 also have identical sequences, with a 42.9% reduction. Alpha–beta therefore preserved minimax decisions while reducing search effort in these matchups. Games 1 and 4 cannot isolate the effect of pruning because they have different starting players.

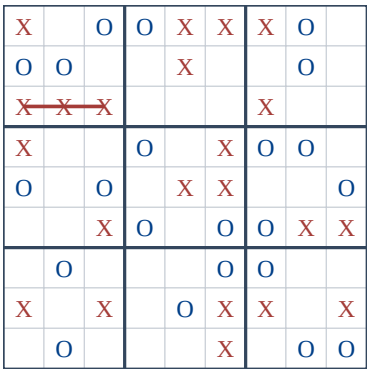
The offensive heuristic emphasizes creating its own threats, whereas the defensive heuristic emphasizes completed blocks. The latter does not directly measure the danger of the board to which the next player is sent, which can leave it vulnerable to sequences of forced routing.



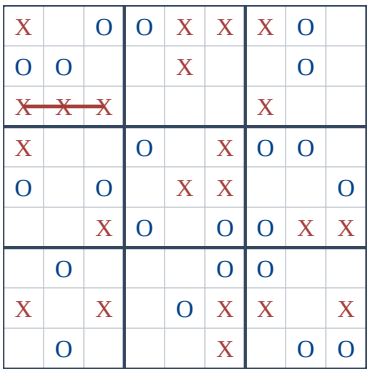
(a) Game 1: X wins; 16,716 nodes



(b) Game 2: X wins; 12,166 nodes



(c) Game 3: X wins; 11,881 nodes



(d) Game 4: X wins; 6,784 nodes

Figure 2: Final boards of the predefined agents.

X is offensive and O is defensive. Thick lines separate local boards; colored lines mark winning triples. The four boards correspond to the rows of Table 2.

Section III Designed Evaluation

Formulation

The designed agent plays O against predefined offensive X. Both use depth-three alpha–beta search. The evaluation considers both players' threats and the board where the next player must actually move. A winning opportunity elsewhere is not necessarily playable on the next turn, so routing is evaluated explicitly.

Let T count unblocked two-in-a-row lines, W sum distinct winning empty cells over local boards, and F count local boards with at least two distinct winning cells. Let C count occupied corners and M occupied centers, with subscripts X and O identifying the player. Set $s = 1$ when X moves next and $s = -1$ when O moves next. Of the next player's L legal moves, S are safe: they do not immediately send the opponent to a board where that opponent can win. When a move routes back to the same local board, its cell is updated before testing the danger.

$$\begin{aligned} H = {} & 400(T_X - T_O) + 35(W_X - W_O) + 800(F_X - F_O) \\ & + 15(C_X - C_O) + 40(M_X - M_O) + 6000s\frac{S}{L}. \end{aligned}$$

Actual terminal states take priority and return ± 10000. If the next player can win immediately, return $9500s$. If every legal move allows an immediate opponent win, return $-9000s$. Otherwise, use H clipped to $[-8000, 8000]$. A full global board with no winner returns zero. These static tactical checks do not extend the recursive search beyond three plies.

The ratio S/L avoids favoring a board merely because it has fewer empty cells. The feature F distinguishes two independent winning cells from two lines sharing one blocking cell. Immediate threats take priority over general positional scores; centers and corners have smaller weights.

Feature interpretation and the search horizon

T measures potential attacking lines; W and F distinguish threats that can be removed by a single block from independent winning cells. Two lines sharing one winning square can still be blocked together. Even independent threats depend on the forced board, so global line counts alone are insufficient.

The safety term has the current player's sign. A larger S/L increases X utility when X moves next and decreases it when O moves next. It converts the consequence of sending the opponent to a particular board into a comparable state feature. Separating actual outcomes, immediate tactical values, and ordinary heuristic values prevents distant line or corner scores from outweighing immediate wins and losses. These scores rank positions; they are not winning probabilities.

Finite depth creates a horizon effect: longer forced returns can fall beyond the search boundary. The evaluation recognizes immediate wins and positions where every choice loses immediately, but it cannot prove the safety of longer sequences. Match results are therefore necessary, and success against the fixed opponent does not establish strength against arbitrary human strategies.

Parameter comparison

Development covered all nine starting boards and both starting players, giving eighteen conditions. The initial safe-count scheme won eleven; the best adjusted safe-count scheme won fourteen. Position weights, threat spread, and safety ratios were then compared. The ratio weight 6000 was selected for the most development wins in Table 3.

Table 3: Development comparison of mobility features.

Mobility feature	Weight	Development wins / 18
Safe-move count	35	11
Safe-move count	1000	14
Safe-move ratio	1000	16
Safe-move ratio	3000	16
Safe-move ratio	6000	18

There are only eighteen opening conditions, so development and subsequent random games are not strictly independent. The results establish an improvement against this specified opponent, rather than a guarantee against arbitrary strategies.

Randomized match results

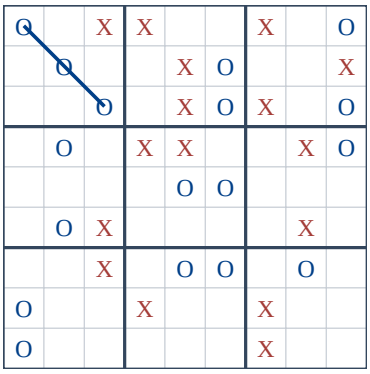
With random seed 448, each game's starting board and player were sampled randomly. Across 40 games, O won 40, X won 0, and 0 were drawn. The designed agent's win rate was 100.0%, exceeding the required 80%. The average total search count was 6,516.7 nodes per game.

Table 4: Per-game results; node counts sum both players' searches.

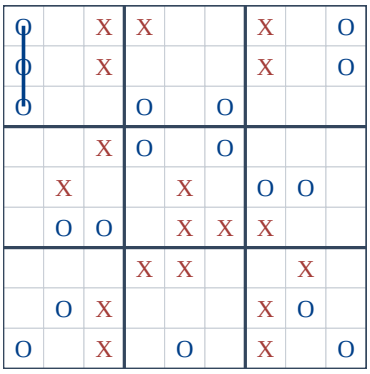
Game	Start	First	Plies	Nodes	Win	Game	Start	First	Plies	Nodes	Win
1	4	O	33	5,759	O	21	3	X	38	6,716	O
2	1	X	36	7,901	O	22	8	O	27	5,217	O
3	4	O	33	5,759	O	23	3	O	29	5,976	O
4	0	X	32	6,700	O	24	8	O	27	5,217	O
5	8	O	27	5,217	O	25	4	X	34	6,104	O
6	6	X	36	7,107	O	26	6	X	36	7,107	O
7	1	X	36	7,901	O	27	6	O	37	7,187	O
8	3	X	38	6,716	O	28	7	X	36	6,206	O
9	8	O	27	5,217	O	29	6	O	37	7,187	O
10	5	X	36	6,536	O	30	1	O	33	6,850	O
11	5	O	41	6,969	O	31	6	O	37	7,187	O
12	7	X	36	6,206	O	32	5	X	36	6,536	O
13	4	X	34	6,104	O	33	3	O	29	5,976	O
14	8	O	27	5,217	O	34	0	X	32	6,700	O
15	5	O	41	6,969	O	35	6	X	36	7,107	O
16	5	X	36	6,536	O	36	3	X	38	6,716	O
17	1	X	36	7,901	O	37	0	O	35	7,409	O
18	0	X	32	6,700	O	38	8	O	27	5,217	O
19	0	O	35	7,409	O	39	3	O	29	5,976	O
20	5	X	36	6,536	O	40	3	X	38	6,716	O

Policies and tie-breaking are deterministic. Repeated starting conditions therefore produce repeated games; the forty runs are not forty distinct tests of opponent strategy.

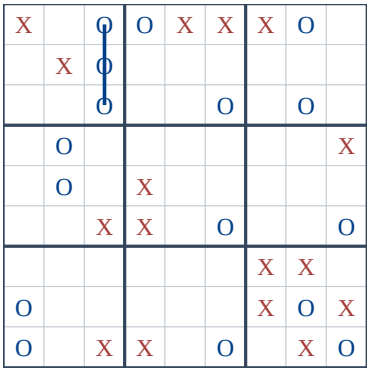
Representative positions



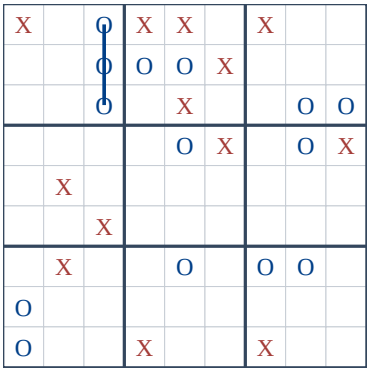
(a) Game 1: start 4, O first



(b) Game 2: start 1, X first



(c) Game 4: start 0, X first



(d) Game 5: start 8, O first

Figure 3: Representative final boards of the designed agent.

Immediately before the final move, X had 5, 7, 3, 4 legal choices in the four games, respectively, but none could prevent O from winning on the next turn. The preceding threats and routing had already forced the outcome; these wins did not merely depend on X overlooking a final block. The predefined offensive agent emphasizes its own two-in-a-row score, whereas the designed evaluation also considers the destination of the next turn and gives routing threats greater priority. Final boards alone cannot isolate the causal effect of an individual weight. These findings remain specific to the tested opponent and opening conditions.

Section IV Human Games

The human plays X and the designed agent plays O, with randomized starting board and player. At least ten games are required, followed by a discussion of win rates, representative final boards, and the agent's strengths and weaknesses. The human experiment has not yet been completed. Its win rate, representative boards, and empirical discussion remain to be added.

Depth-three search can miss longer forced-routing sequences. The safety ratio checks immediate danger, not safety across several future turns. Human games should examine whether threats prepared on several boards can force the agent into an unavoidable loss, and whether fixed corner and center weights interfere with necessary defense. These are questions for the planned experiment, not empirical conclusions from human play.

Extra Credit

The global three-board-in-a-row extension was not implemented.

Statement of Contribution

The following is a proposed allocation for distributing the remaining work, not a record of completed individual contributions. The group should review the results together and update this statement after completing the assigned tasks.

Table 5: Proposed allocation, pending group confirmation.

Member	Proposed responsibility
Wang Libin	CSP formulation, tiling solver review, correctness checks, and Section I.
Jiang Jiahao	Minimax and alpha–beta review, four predefined matchups, node statistics, and Section II.
Wang Zizheng	Designed evaluation, parameter comparison, randomized game analysis, and Section III.
Xue Xiaokang	Organizing at least ten human games, recording representative boards, Section IV, and coordinating proofreading.

The current code, automated experiments, and report draft were prepared with Codex assistance.